



Priority Queues

CSE 332 21AU
Lecture 3

Formally Big-O

We wanted to find an upper bound on our algorithm's running time, but

- We don't want to care about constant factors.
- We only care about what happens as n gets large.

The formal, mathematical definition is Big-O.

Big-O

$f(n)$ is $O(g(n))$ if there exist positive constants c, n_0 such that for all $n \geq n_0$,

$$f(n) \leq c \cdot g(n)$$

We also say that $g(n)$ "dominates" $f(n)$.

Writing Proofs

Claim: For every odd integer y , there exists an even integer x , such that $x > y$.

Proof:

Let y be an arbitrary odd integer. By definition, $y = 2z + 1$ for some integer z .

Consider $x = 2(z + 1)$.

x is even (since it can be written as 2 times some integer) and

$x = 2(z + 1) = 2z + 2 > 2z + 1 = y$. So x meets both of the required properties. \square

Writing Proofs

Where did that $x = 2(z + 1)$ come from?

You probably came up with that even integer first, before you started writing the proof.

That was some “scratch work” – the insight isn’t explained in the final proof

- You just say “Consider”

Don’t try to skip the scratch work when drafting your big-O proofs.

- But it won’t appear in your final version.

Using the Definition

Let's show: $10n^2 + 15n$ is $O(n^2)$

Using the Definition

Let's show: $10n^2 + 15n$ is $O(n^2)$

Recreation of whiteboard:

Scratch work:

$$10n^2 \leq 10n^2$$

$$15n \leq 15n^2 \text{ for } n \geq 1$$

$$10n^2 + 15n \leq 25n^2 \text{ for } n \geq 1$$

Proof:


Take $c = 25$ and $n_0 = 1$.

The inequality $10n^2 \leq 10n^2$ is always true. The inequality $15n \leq 15n^2$ is true for $n \geq 1$, as the right hand side is a factor of n more than the right hand side.

As long as both inequalities are true we can add them, thus

$$10n^2 + 15n \leq 25n^2 \text{ holds as long as } n \geq 1.$$

This is exactly the inequality we needed to show.



A new ADT: Priority Queues

A New ADT

Our previous worklists (stacks, queues, etc.) all choose the next element based on the order they were inserted.

That's not always a good idea.

Emergency rooms aren't first-come-first-served.

Sometimes our objects come with a **priority**, that tells us what we need to do next.

An ADT that can handle a line with priorities is a **priority queue**.

Priority Queue ADT

Min Priority Queue ADT

state

Set of comparable values
- Ordered based on "priority"

behavior

insert(value) – add a new element to the collection

removeMin() – returns the element with the smallest priority, removes it from the collection

peekMin() – find, but do not remove the element with the smallest priority

Uses:

- Operating System
- Well-designed printers
- Huffman Codes (in 143)
- Sorting (in Project 2)
- Graph algorithms

Implementing Priority Queues: Take I

Maybe we already know how to implement a priority queue.
How long would insert and removeMin take with these data structures?

	Insert	removeMin
Unsorted Array		
Unsorted Linked List		
Sorted Linked List		
Sorted Circular Array		
Binary Search Tree		




For Array implementations, assume that the array is not yet full.
Other than this assumption, do **worst case** analysis.

Review: Binary Search Trees

A BST is:

1. A binary tree
2. For each node, everything in its left subtree is smaller than it and everything in its right subtree is larger than it.

Are These BSTs?



Implementing Priority Queues: Take I

Maybe we already know how to implement a priority queue.
How long would `insert` and `removeMin` take with these data structures?

	Insert	removeMin
Unsorted Array	$\Theta(1)$	$\Theta(n)$
Unsorted Linked List	$\Theta(1)$	$\Theta(n)$
Sorted Linked List	$\Theta(n)$	$\Theta(1)$
Sorted Circular Array	$\Theta(n)$	$\Theta(1)$
Binary Search Tree		

For Array implementations, assume that the array is not yet full.
Other than this assumption, do **worst case** analysis.

Implementing Priority Queues: Take I

Maybe we already know how to implement a priority queue.
How long would `insert` and `removeMin` take with these data structures?

	Insert	removeMin
Unsorted Array	$\Theta(1)$	$\Theta(n)$
Unsorted Linked List	$\Theta(1)$	$\Theta(n)$
Sorted Linked List	$\Theta(n)$	$\Theta(1)$
Sorted Circular Array	$\Theta(n)$	$\Theta(1)$
Binary Search Tree	$\Theta(n)$	$\Theta(n)$

For Array implementations, assume that the array is not yet full.
Other than this assumption, do **worst case** analysis.

Implementing Priority Queues: Take II

BSTs have really bad behavior in the **worst** case, but is it actually a common problem?

Fact: On average, the height of a BST is $O(\log n)$
(for some suitable definition of “average”)

Can we somehow get that behavior in the **worst case** for priority queues?

BST Properties

A BST is:

1. A binary tree
2. For each node, everything in its left subtree is smaller than it and everything in its right subtree is larger than it.

Point 2 is what causes the really bad behavior in the worst case.

We probably don't want exactly that requirement for implementing a priority queue.

Maybe we can explicitly enforce that we don't get a degenerate tree.



Binary Heaps

A Binary Heap is

1. A Binary Tree
2. Every node is less than or equal to all of its children
 - In particular, the smallest element must be the root!
3. The tree is **complete**
 - Every level of the tree is completely filled, except possibly the last row, which is filled from left to right.
 - No degenerate trees!

Tree Words


Height – the number of edges on the longest path from the root to a leaf.




Tree Words

Complete – every row is completely filled, except possibly the last row, which is filled from left to right.


Perfect – every row is completely filled



Complete, but not perfect



Neither



Both Perfect and Complete

Heights of Perfect Trees

How many nodes are there in level i of a perfect binary tree?

Heights of Perfect Trees

How many nodes are there in level i of a perfect binary tree?

On the whiteboard we derived that the number of nodes on level i of a binary tree was 2^i .

Thus the total number of nodes in a perfect binary tree of height h is

$$\sum_{i=0}^h 2^i = 2^{h+1} - 1.$$

So if we have n nodes in a perfect tree, we can use the formula

$$n = 2^{h+1} - 1 \text{ to conclude that } h = O(\log n), \text{ so}$$

A perfect tree with n nodes has height $O(\log n)$.

A similar argument can show the same statement for complete trees.


Binary Heaps

A Binary Heap is

1. A Binary Tree
2. Every node is less than or equal to all of its children
 - In particular, the smallest element must be the root!
3. The tree is **complete**
 - Every level of the tree is completely filled, except possibly the last row, which is filled from left to right.
 - No degenerate trees!

Implementing Heaps

Let's start with removeMin.




Idea: take the bottom right-most node and use it to plug the hole

Shape is correct now

But that value might be too big. We need to "percolate it down"

Implementing Heaps

Insertion



What is the shape going to be after the insertion?

Again, plug the hole first.

Might violate the heap property. Percolate it up



An Optimization

Pointers are annoying.

They're also slow.

Our tree is so nicely shaped, we don't need pointers

We can use an array instead.



An Optimization


If I'm at index i , what is the index of:

My left child, right child and parent?

My left child: $2i$

My right child: $2i + 1$

My parent: $\left\lfloor \frac{i}{2} \right\rfloor$



	1	4	3	5	6	8	7			
0	1	2	3	4	5	6	7	8	9	10